

Keras -- MLPs on MNIST

In [1]:

```
# if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow" use this command
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
```

Using TensorFlow backend.

In [2]:

```
%matplotlib notebook
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

In [3]:

```
# the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

In [4]:

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d, %d)"%(X_train.shape[1], X_train.shape[2]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d, %d)"%(X_test.shape[1], X_test.shape[2]))
```

Number of training examples : 60000 and each image is of shape (28, 28)
Number of training examples : 10000 and each image is of shape (28, 28)

In [5]:

```
# if you observe the input shape its 2 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784

X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

In [6]:

```
# after converting the input images from 3d to 2d vectors

print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d)"%(X_train.shape[1]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d)"%(X_test.shape[1]))
```

Number of training examples : 60000 and each image is of shape (784)
Number of training examples : 10000 and each image is of shape (784)

```
# An example data point
print(X_train[0])
```

In [8]:

```
# if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms lets try to normalize the data
#  $X \Rightarrow (X - X_{min}) / (X_{max} - X_{min}) = X / 255$ 

X_train = X_train/255
X_test = X_test/255
```

```
# example data point after normalizing
print(X_train[0])
```

[illegible]

[illegible]

[illegible]

In [10]:

```
# here we are having a class number for each image
print("Class label of first image :", y_train[0])

# lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# this conversion needed for MLPs

Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ",Y_train[0])
```

```
Class label of first image : 5
After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]
```

Softmax classifier

In [11]:

```
#importing
from keras.models import Sequential
from keras.layers import Dense, Activation
from keras.layers.normalization import BatchNormalization
from keras.layers import Dropout
from prettytable import PrettyTable
```

In [12]:

```
# some model parameters

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20
```

1.1)MLP with 2 hidden layers + Relu + BatchNormalisation +Dropout+ AdamOptimizer

In [14]:

```
model_relu = Sequential()
#Let's take dimension of first and second hidden layer 400 and 250 respectively
# If we sample weights from a normal distribution  $N(0,\sigma)$  we satisfy this condition with  $\sigma=\sqrt{2/(n_i)}$ .
# h1 =>  $\sigma=\sqrt{2/(fan\_in)} = 0.070 \Rightarrow N(0,\sigma) = N(0,0.070)$ 
# h2 =>  $\sigma=\sqrt{2/(fan\_in)} = 0.089 \Rightarrow N(0,\sigma) = N(0,0.089)$ 

model_relu.add(Dense(400, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.070, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))
model_relu.add(Dense(250, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.089, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))
model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()
```

Layer (type)	Output Shape	Param #
dense_4 (Dense)	(None, 400)	314000
batch_normalization_3 (Batch Normalization)	(None, 400)	1600
dropout_3 (Dropout)	(None, 400)	0
dense_5 (Dense)	(None, 250)	100250
batch_normalization_4 (Batch Normalization)	(None, 250)	1000
dropout_4 (Dropout)	(None, 250)	0
dense_6 (Dense)	(None, 10)	2510
Total params: 419,360		
Trainable params: 418,060		
Non-trainable params: 1,300		

In [15]:

```
#Applying adam optimizer
model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

W0804 15:39:09.342900 6892 deprecation_wrapper.py:119] From C:\Users\Hi\Anaconda3\lib\site-packages\keras\optimizers.py:790: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.

W0804 15:39:10.184588 6892 deprecation.py:323] From C:\Users\Hi\Anaconda3\lib\site-packages\tensorflow\python\ops\math_grad.py:1250: add_dispatch_support.<locals>.wrapper (from tensorflow.python.ops.array_ops) is deprecated and will be removed in a future version.

Instructions for updating:
Use tf.where in 2.0, which has the same broadcast rule as np.where

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 34s 572us/step - loss: 0.4713 - acc: 0.8576 - val_loss: 0.1440 - val_acc: 0.9551
Epoch 2/20
60000/60000 [=====] - 26s 427us/step - loss: 0.2234 - acc: 0.9323 - val_loss: 0.1106 - val_acc: 0.9661
Epoch 3/20
60000/60000 [=====] - 26s 430us/step - loss: 0.1717 - acc: 0.9467 - val_loss: 0.0978 - val_acc: 0.9700
Epoch 4/20
60000/60000 [=====] - 25s 424us/step - loss: 0.1457 - acc: 0.9557 - val_loss: 0.0834 - val_acc: 0.9739
Epoch 5/20
60000/60000 [=====] - 25s 424us/step - loss: 0.1289 - acc: 0.9608 - val_loss: 0.0761 - val_acc: 0.9760
Epoch 6/20
60000/60000 [=====] - 25s 423us/step - loss: 0.1160 - acc: 0.9641 - val_loss: 0.0728 - val_acc: 0.9768
Epoch 7/20
60000/60000 [=====] - 23s 383us/step - loss: 0.1060 - acc: 0.9669 - val_loss: 0.0684 - val_acc: 0.9792
Epoch 8/20
60000/60000 [=====] - 21s 348us/step - loss: 0.0953 - acc: 0.9696 - val_loss: 0.0629 - val_acc: 0.9792
Epoch 9/20
60000/60000 [=====] - 26s 426us/step - loss: 0.0932 - acc: 0.9709 - val_loss: 0.0614 - val_acc: 0.9811
Epoch 10/20
60000/60000 [=====] - 25s 408us/step - loss: 0.0855 - acc: 0.9730 - val_loss: 0.0650 - val_acc: 0.9802
Epoch 11/20
60000/60000 [=====] - 26s 427us/step - loss: 0.0809 - acc: 0.9744 - val_loss: 0.0625 - val_acc: 0.9807
Epoch 12/20
60000/60000 [=====] - 25s 425us/step - loss: 0.0775 - acc: 0.9753 - val_loss: 0.0615 - val_acc: 0.9813
Epoch 13/20
60000/60000 [=====] - 25s 419us/step - loss: 0.0729 - acc: 0.9766 - val_loss: 0.0596 - val_acc: 0.9814
Epoch 14/20
60000/60000 [=====] - 25s 420us/step - loss: 0.0690 - acc: 0.9775 - val_loss: 0.0581 - val_acc: 0.9817
Epoch 15/20
60000/60000 [=====] - 25s 417us/step - loss: 0.0678 - acc: 0.9786 - val_loss: 0.0576 - val_acc: 0.9821
Epoch 16/20
60000/60000 [=====] - 25s 421us/step - loss: 0.0635 - acc: 0.9796 - val_loss: 0.0610 - val_acc: 0.9821
Epoch 17/20
60000/60000 [=====] - 26s 427us/step - loss: 0.0611 - acc: 0.9803 - val_loss: 0.0599 - val_acc: 0.9831
Epoch 18/20
60000/60000 [=====] - 25s 416us/step - loss: 0.0594 - acc: 0.9808 - val_loss: 0.0568 - val_acc: 0.9836
Epoch 19/20
60000/60000 [=====] - 26s 437us/step - loss: 0.0535 - acc: 0.9824 - val_loss: 0.0598 - val_acc: 0.9825
Epoch 20/20
60000/60000 [=====] - 27s 453us/step - loss: 0.0551 - acc: 0.9819 - val_loss: 0.0554 - val_acc: 0.9839
```

In [16]:

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])


fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))
```

```
x = list(range(1, nb_epoch+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.05540496020909195
Test accuracy: 0.9839



In [17]:


```
w_after = model_relu.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```





1.2)MLP with 2 hidden layers + Relu + AdamOptimizer(WithoutBatchNormalisation,Without dropout)

In [25]:

```
model_relu = Sequential()
#Let's take dimension of first and second hidden layer 400 and 250 respectively
# If we sample weights from a normal distribution  $N(0,\sigma)$  we satisfy this condition with  $\sigma=\sqrt{2/(n_i)}$ .
# h1 =>  $\sigma=\sqrt{2/(fan\_in)} = 0.070 \Rightarrow N(0,\sigma) = N(0,0.070)$ 
# h2 =>  $\sigma=\sqrt{2/(fan\_in)} = 0.089 \Rightarrow N(0,\sigma) = N(0,0.089)$ 

model_relu.add(Dense(400, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.070, seed=None)))
model_relu.add(Dense(250, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.089, seed=None)))
model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()
```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_14 (Dense)	(None, 400)	314000
dense_15 (Dense)	(None, 250)	100250
dense_16 (Dense)	(None, 10)	2510
=====	=====	=====
Total params: 416,760		
Trainable params: 416,760		
Non-trainable params: 0		

In [26]:

```
#Applying adam optimizer
model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 9s 150us/step - loss: 0.2348 - acc: 0.9305 - val_loss: 0.1079 - val_acc: 0.9656
Epoch 2/20
60000/60000 [=====] - 9s 146us/step - loss: 0.0862 - acc: 0.9738 - val_loss: 0.0815 - val_acc: 0.9747
Epoch 3/20
60000/60000 [=====] - 9s 153us/step - loss: 0.0524 - acc: 0.9835 - val_loss: 0.0708 - val_acc: 0.9776
```



```

val_loss: 0.0574 - val_acc: 0.9824
Epoch 4/20
60000/60000 [=====] - 9s 156us/step - loss: 0.0353 - acc: 0.9889 -
val_loss: 0.0574 - val_acc: 0.9824
Epoch 5/20
60000/60000 [=====] - 9s 151us/step - loss: 0.0268 - acc: 0.9908 -
val_loss: 0.0736 - val_acc: 0.9772
Epoch 6/20
60000/60000 [=====] - 9s 153us/step - loss: 0.0190 - acc: 0.9940 -
val_loss: 0.0690 - val_acc: 0.9804
Epoch 7/20
60000/60000 [=====] - 9s 152us/step - loss: 0.0156 - acc: 0.9949 -
val_loss: 0.0756 - val_acc: 0.9768
Epoch 8/20
60000/60000 [=====] - 9s 156us/step - loss: 0.0186 - acc: 0.9935 -
val_loss: 0.0655 - val_acc: 0.9808
Epoch 9/20
60000/60000 [=====] - 9s 153us/step - loss: 0.0111 - acc: 0.9963 -
val_loss: 0.0814 - val_acc: 0.9782
Epoch 10/20
60000/60000 [=====] - 10s 159us/step - loss: 0.0096 - acc: 0.9971 - val_l
oss: 0.0904 - val_acc: 0.9786
Epoch 11/20
60000/60000 [=====] - 9s 155us/step - loss: 0.0155 - acc: 0.9946 -
val_loss: 0.0791 - val_acc: 0.9798
Epoch 12/20
60000/60000 [=====] - 9s 156us/step - loss: 0.0098 - acc: 0.9966 -
val_loss: 0.0842 - val_acc: 0.9789
Epoch 13/20
60000/60000 [=====] - 9s 152us/step - loss: 0.0110 - acc: 0.9964 -
val_loss: 0.0717 - val_acc: 0.9821
Epoch 14/20
60000/60000 [=====] - 9s 155us/step - loss: 0.0074 - acc: 0.9977 -
val_loss: 0.0895 - val_acc: 0.9803
Epoch 15/20
60000/60000 [=====] - 10s 172us/step - loss: 0.0065 - acc: 0.9979 - val_l
oss: 0.0744 - val_acc: 0.9832
Epoch 16/20
60000/60000 [=====] - 12s 194us/step - loss: 0.0110 - acc: 0.9965 - val_l
oss: 0.1034 - val_acc: 0.9791
Epoch 17/20
60000/60000 [=====] - 10s 161us/step - loss: 0.0122 - acc: 0.9958 - val_l
oss: 0.0782 - val_acc: 0.9802
Epoch 18/20
60000/60000 [=====] - 10s 165us/step - loss: 0.0034 - acc: 0.9991 - val_l
oss: 0.0870 - val_acc: 0.9821
Epoch 19/20
60000/60000 [=====] - 10s 173us/step - loss: 0.0074 - acc: 0.9973 - val_l
oss: 0.0937 - val_acc: 0.9791
Epoch 20/20
60000/60000 [=====] - 11s 186us/step - loss: 0.0073 - acc: 0.9975 - val_l
oss: 0.0829 - val_acc: 0.9835

```

In [27]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])


fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.08285639366751603
Test accuracy: 0.9835



In [28]:


```
w_after = model_relu.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```





2.1)MLP with 3 hidden layers + Relu + BatchNormalisation+Dropout + AdamOptimizer

In [18]:

```
model_relu = Sequential()
#Let's take dimension of first, second and third hidden layer 480, 280 and 80 respectively
# If we sample weights from a normal distribution  $N(0,\sigma)$  we satisfy this condition with
 $\sigma = \sqrt{2/(fan\_in)}$ .
# h1 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.064 \Rightarrow N(0,\sigma) = N(0,0.064)$ 
# h2 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.084 \Rightarrow N(0,\sigma) = N(0,0.084)$ 
# h3 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.158 \Rightarrow N(0,\sigma) = N(0,0.158)$ 

model_relu.add(Dense(480, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.064, seed=None)))
model_relu.add(Dropout(0.5))
model_relu.add(BatchNormalization())

model_relu.add(Dense(280, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.084, seed=None)))
model_relu.add(Dropout(0.5))

model_relu.add(Dense(80, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.158, seed=None)))
model_relu.add(Dropout(0.5))
model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()
```

Layer (type)	Output Shape	Param #
dense_7 (Dense)	(None, 480)	376800
dropout_5 (Dropout)	(None, 480)	0
batch_normalization_5 (Batch Normalization)	(None, 480)	1920
dense_8 (Dense)	(None, 280)	134680
dropout_6 (Dropout)	(None, 280)	0
dense_9 (Dense)	(None, 80)	22480
dropout_7 (Dropout)	(None, 80)	0
dense_10 (Dense)	(None, 10)	810
Total params: 536,690		
Trainable params: 535,730		
Non-trainable params: 960		

In [19]:

```
#Applying adam optimizer
model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 24s 405us/step - loss: 1.3775 - acc: 0.6410 - val_loss: 0.2898 - val_acc: 0.9212

Epoch 2/20

```

60000/60000 [=====] - 31s 517us/step - loss: 0.5014 - acc: 0.8528 - val_1
oss: 0.1975 - val_acc: 0.9427
Epoch 3/20
60000/60000 [=====] - 35s 587us/step - loss: 0.3625 - acc: 0.8976 - val_1
oss: 0.1517 - val_acc: 0.9554
Epoch 4/20
60000/60000 [=====] - 31s 513us/step - loss: 0.2964 - acc: 0.9192 - val_1
oss: 0.1333 - val_acc: 0.9603
Epoch 5/20
60000/60000 [=====] - 32s 528us/step - loss: 0.2493 - acc: 0.9306 - val_1
oss: 0.1180 - val_acc: 0.9636
Epoch 6/20
60000/60000 [=====] - 33s 553us/step - loss: 0.2262 - acc: 0.9384 - val_1
oss: 0.1123 - val_acc: 0.9674
Epoch 7/20
60000/60000 [=====] - 32s 536us/step - loss: 0.2018 - acc: 0.9448 - val_1
oss: 0.1069 - val_acc: 0.9704
Epoch 8/20
60000/60000 [=====] - 29s 488us/step - loss: 0.1898 - acc: 0.9481 - val_1
oss: 0.1008 - val_acc: 0.9716
Epoch 9/20
60000/60000 [=====] - 31s 524us/step - loss: 0.1771 - acc: 0.9525 - val_1
oss: 0.0951 - val_acc: 0.9729
Epoch 10/20
60000/60000 [=====] - 31s 522us/step - loss: 0.1630 - acc: 0.9562 - val_1
oss: 0.0943 - val_acc: 0.9740
Epoch 11/20
60000/60000 [=====] - 34s 565us/step - loss: 0.1552 - acc: 0.9577 - val_1
oss: 0.1048 - val_acc: 0.9727
Epoch 12/20
60000/60000 [=====] - 32s 527us/step - loss: 0.1483 - acc: 0.9602 - val_1
oss: 0.0912 - val_acc: 0.9750
Epoch 13/20
60000/60000 [=====] - 34s 559us/step - loss: 0.1397 - acc: 0.9620 - val_1
oss: 0.0821 - val_acc: 0.9766
Epoch 14/20
60000/60000 [=====] - 31s 514us/step - loss: 0.1369 - acc: 0.9625 - val_1
oss: 0.0909 - val_acc: 0.9764
Epoch 15/20
60000/60000 [=====] - 30s 505us/step - loss: 0.1251 - acc: 0.9641 - val_1
oss: 0.0833 - val_acc: 0.9769
Epoch 16/20
60000/60000 [=====] - 31s 514us/step - loss: 0.1261 - acc: 0.9653 - val_1
oss: 0.0772 - val_acc: 0.9779
Epoch 17/20
60000/60000 [=====] - 26s 440us/step - loss: 0.1258 - acc: 0.9661 - val_1
oss: 0.0764 - val_acc: 0.9779
Epoch 18/20
60000/60000 [=====] - 17s 284us/step - loss: 0.1116 - acc: 0.9688 - val_1
oss: 0.0743 - val_acc: 0.9796
Epoch 19/20
60000/60000 [=====] - 18s 303us/step - loss: 0.1157 - acc: 0.9679 - val_1
oss: 0.0762 - val_acc: 0.9798
Epoch 20/20
60000/60000 [=====] - 14s 225us/step - loss: 0.1058 - acc: 0.9701 - val_1
oss: 0.0738 - val_acc: 0.9809

```

In [20]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

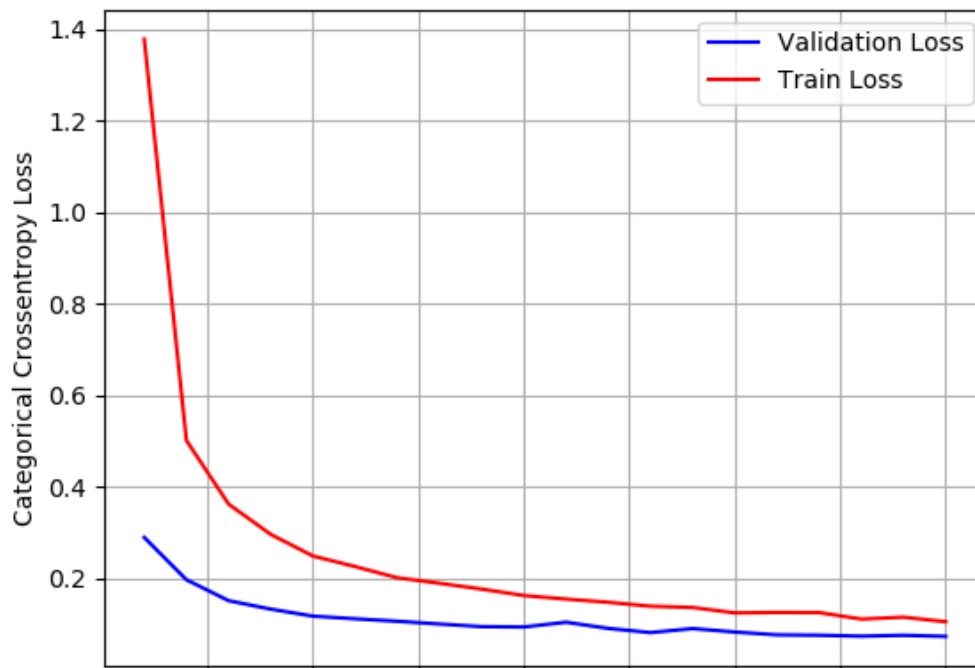
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.0738352784443341
Test accuracy: 0.9809



In [21]:


```
w_after = model_relu.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```





2.2)MLP with 3 hidden layers + Relu + AdamOptimizer(WithoutBatchNormalisation,Without dropout)

In [29]:

```
model_relu = Sequential()
#Let's take dimension of first, second and third hidden layer 480, 280 and 80 respectively
# If we sample weights from a normal distribution  $N(0,\sigma)$  we satisfy this condition with  $\sigma=\sqrt{2/(n_i)}$ .
# h1 =>  $\sigma=\sqrt{2/(fan\_in)} = 0.064 \Rightarrow N(0,\sigma) = N(0,0.064)$ 
# h2 =>  $\sigma=\sqrt{2/(fan\_in)} = 0.084 \Rightarrow N(0,\sigma) = N(0,0.084)$ 
# h3 =>  $\sigma=\sqrt{2/(fan\_in)} = 0.158 \Rightarrow N(0,\sigma) = N(0,0.158)$ 

model_relu.add(Dense(480, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.064, seed=None)))

model_relu.add(Dense(280, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.084, seed=None)))

model_relu.add(Dense(80, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.158, seed=None)))
model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()
```

Layer (type)	Output Shape	Param #
dense_17 (Dense)	(None, 480)	376800
dense_18 (Dense)	(None, 280)	134680
dense_19 (Dense)	(None, 80)	22480
dense_20 (Dense)	(None, 10)	810
Total params: 534,770		
Trainable params: 534,770		
Non-trainable params: 0		

In [30]:

```
#Applying adam optimizer
model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 12s 205us/step - loss: 0.2225 - acc: 0.9318 - val_loss: 0.1375 - val_acc: 0.9572
Epoch 2/20
60000/60000 [=====] - 11s 188us/step - loss: 0.0835 - acc: 0.9741 - val_loss: 0.0897 - val_acc: 0.9718
Epoch 3/20
60000/60000 [=====] - 12s 193us/step - loss: 0.0514 - acc: 0.9833 - val_loss: 0.0896 - val_acc: 0.9724
Epoch 4/20
60000/60000 [=====] - 12s 182us/step - loss: 0.0386 - acc: 0.9873 - val_loss: 0.0386 - val_acc: 0.9873
```

```

60000/60000 [=====] - 12s 192us/step - loss: 0.0366 - acc: 0.9873 - val_loss: 0.0933 - val_acc: 0.9731
Epoch 5/20
60000/60000 [=====] - 12s 206us/step - loss: 0.0288 - acc: 0.9906 - val_loss: 0.0727 - val_acc: 0.9804
Epoch 6/20
60000/60000 [=====] - 11s 189us/step - loss: 0.0218 - acc: 0.9930 - val_loss: 0.0835 - val_acc: 0.9777
Epoch 7/20
60000/60000 [=====] - 11s 191us/step - loss: 0.0216 - acc: 0.9930 - val_loss: 0.0931 - val_acc: 0.9751
Epoch 8/20
60000/60000 [=====] - 12s 194us/step - loss: 0.0192 - acc: 0.9935 - val_loss: 0.0788 - val_acc: 0.9787
Epoch 9/20
60000/60000 [=====] - 11s 189us/step - loss: 0.0157 - acc: 0.9950 - val_loss: 0.0809 - val_acc: 0.9792
Epoch 10/20
60000/60000 [=====] - 11s 189us/step - loss: 0.0185 - acc: 0.9940 - val_loss: 0.0821 - val_acc: 0.9801
Epoch 11/20
60000/60000 [=====] - 11s 190us/step - loss: 0.0120 - acc: 0.9960 - val_loss: 0.0985 - val_acc: 0.9783
Epoch 12/20
60000/60000 [=====] - 11s 184us/step - loss: 0.0161 - acc: 0.9949 - val_loss: 0.0820 - val_acc: 0.9802
Epoch 13/20
60000/60000 [=====] - 11s 191us/step - loss: 0.0117 - acc: 0.9960 - val_loss: 0.0742 - val_acc: 0.9824
Epoch 14/20
60000/60000 [=====] - 11s 187us/step - loss: 0.0119 - acc: 0.9963 - val_loss: 0.0831 - val_acc: 0.9829
Epoch 15/20
60000/60000 [=====] - 12s 197us/step - loss: 0.0133 - acc: 0.9960 - val_loss: 0.1034 - val_acc: 0.9775
Epoch 16/20
60000/60000 [=====] - 12s 200us/step - loss: 0.0122 - acc: 0.9963 - val_loss: 0.0914 - val_acc: 0.9804
Epoch 17/20
60000/60000 [=====] - 12s 202us/step - loss: 0.0097 - acc: 0.9970 - val_loss: 0.0945 - val_acc: 0.9805
Epoch 18/20
60000/60000 [=====] - 12s 201us/step - loss: 0.0094 - acc: 0.9970 - val_loss: 0.0982 - val_acc: 0.9825
Epoch 19/20
60000/60000 [=====] - 13s 223us/step - loss: 0.0099 - acc: 0.9970 - val_loss: 0.1012 - val_acc: 0.9795
Epoch 20/20
60000/60000 [=====] - 12s 202us/step - loss: 0.0108 - acc: 0.9966 - val_loss: 0.0948 - val_acc: 0.9814

```

In [19]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])


fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

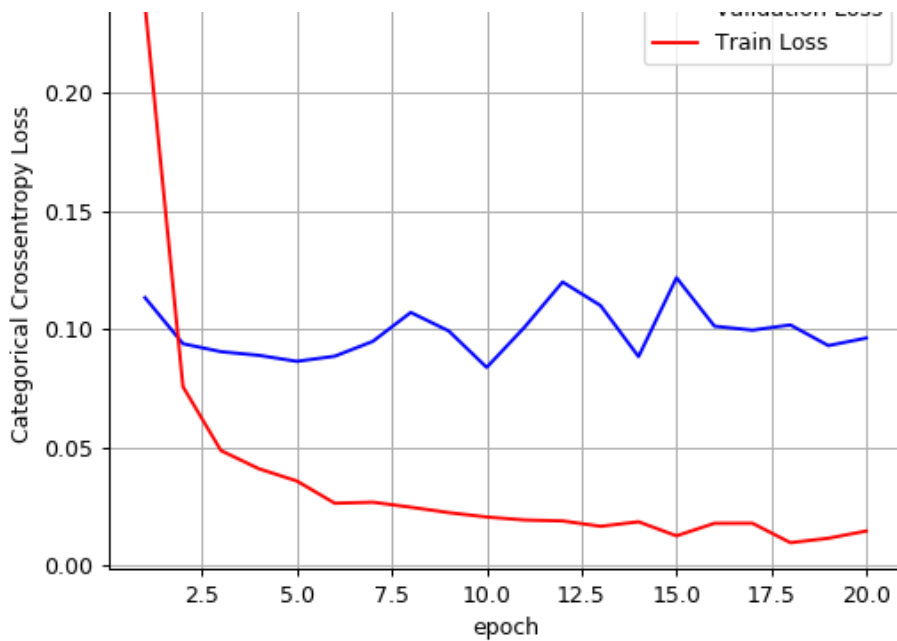
# list of epoch numbers
x = list(range(1,nb_epoch+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.09616601888963419
Test accuracy: 0.9786





In [20]:


```
w_after = model_relu.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```





3.1)MLP with 5 hidden layers + Relu + BatchNormalisation +Dropout+ AdamOptimizer

In [22]:

```
model_relu = Sequential()
#Let's take dimension of first, second,third,fourth and fifth hidden layer 650,550,400,200 and 100
respectively
# If we sample weights from a normal distribution  $N(0,\sigma)$  we satisfy this condition with
 $\sigma = \sqrt{2/(n_i)}$ .
# h1 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.055 \Rightarrow N(0,\sigma) = N(0,0.055)$ 
# h2 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.060 \Rightarrow N(0,\sigma) = N(0,0.060)$ 
# h3 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.070 \Rightarrow N(0,\sigma) = N(0,0.070)$ 
# h4 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.1 \Rightarrow N(0,\sigma) = N(0,0.1)$ 
# h5 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.141 \Rightarrow N(0,\sigma) = N(0,0.141)$ 

model_relu.add(Dense(650, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.055, seed=None)))
model_relu.add(Dropout(0.5))

model_relu.add(Dense(550, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.060, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))

model_relu.add(Dense(400, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.070, seed=None)))
model_relu.add(Dropout(0.5))

model_relu.add(Dense(200, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.1, seed=None)))
model_relu.add(BatchNormalization())
model_relu.add(Dropout(0.5))

model_relu.add(Dense(100, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.141, seed=None)))
model_relu.add(Dropout(0.5))

model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()
```

Layer (type)	Output Shape	Param #
dense_11 (Dense)	(None, 650)	510250
dropout_8 (Dropout)	(None, 650)	0
dense_12 (Dense)	(None, 550)	358050
batch_normalization_6 (Batch Normalization)	(None, 550)	2200
dropout_9 (Dropout)	(None, 550)	0
dense_13 (Dense)	(None, 400)	220400
dropout_10 (Dropout)	(None, 400)	0
dense_14 (Dense)	(None, 200)	80200
batch_normalization_7 (Batch Normalization)	(None, 200)	800
dropout_11 (Dropout)	(None, 200)	0
dense_15 (Dense)	(None, 100)	20100
dropout_12 (Dropout)	(None, 100)	0

```
dense_16 (Dense)                (None, 10)                1010
```

```
=====
Total params: 1,193,010
Trainable params: 1,191,510
Non-trainable params: 1,500
```

In [23]:

```
#Applying adam optimizer
model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 44s 731us/step - loss: 1.6782 - acc: 0.5578 - val_loss: 0.3243 - val_acc: 0.9075
Epoch 2/20
60000/60000 [=====] - 40s 668us/step - loss: 0.4796 - acc: 0.8577 - val_loss: 0.1773 - val_acc: 0.9468
Epoch 3/20
60000/60000 [=====] - 30s 493us/step - loss: 0.3129 - acc: 0.9120 - val_loss: 0.1405 - val_acc: 0.9598
Epoch 4/20
60000/60000 [=====] - 41s 682us/step - loss: 0.2509 - acc: 0.9321 - val_loss: 0.1210 - val_acc: 0.9663
Epoch 5/20
60000/60000 [=====] - 35s 579us/step - loss: 0.2115 - acc: 0.9434 - val_loss: 0.1080 - val_acc: 0.9698
Epoch 6/20
60000/60000 [=====] - 40s 673us/step - loss: 0.1874 - acc: 0.9506 - val_loss: 0.1018 - val_acc: 0.9710
Epoch 7/20
60000/60000 [=====] - 46s 767us/step - loss: 0.1670 - acc: 0.9552 - val_loss: 0.1011 - val_acc: 0.9731
Epoch 8/20
60000/60000 [=====] - 44s 727us/step - loss: 0.1559 - acc: 0.9591 - val_loss: 0.0960 - val_acc: 0.9733
Epoch 9/20
60000/60000 [=====] - 45s 748us/step - loss: 0.1363 - acc: 0.9638 - val_loss: 0.0863 - val_acc: 0.9774
Epoch 10/20
60000/60000 [=====] - 36s 593us/step - loss: 0.1348 - acc: 0.9644 - val_loss: 0.0803 - val_acc: 0.9780
Epoch 11/20
60000/60000 [=====] - 36s 593us/step - loss: 0.1247 - acc: 0.9671 - val_loss: 0.0776 - val_acc: 0.9789
Epoch 12/20
60000/60000 [=====] - 46s 768us/step - loss: 0.1171 - acc: 0.9681 - val_loss: 0.0723 - val_acc: 0.9805
Epoch 13/20
60000/60000 [=====] - 42s 707us/step - loss: 0.1069 - acc: 0.9714 - val_loss: 0.0793 - val_acc: 0.9794
Epoch 14/20
60000/60000 [=====] - 37s 621us/step - loss: 0.1078 - acc: 0.9724 - val_loss: 0.0683 - val_acc: 0.9827
Epoch 15/20
60000/60000 [=====] - 34s 568us/step - loss: 0.0970 - acc: 0.9746 - val_loss: 0.0725 - val_acc: 0.9820
Epoch 16/20
60000/60000 [=====] - 33s 557us/step - loss: 0.0920 - acc: 0.9749 - val_loss: 0.0698 - val_acc: 0.9823
Epoch 17/20
60000/60000 [=====] - 37s 618us/step - loss: 0.0776 - acc: 0.9787 - val_loss: 0.0704 - val_acc: 0.9829
```

In [24]:

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('Epoch')
ax.set_ylabel('Categorical Crossentropy Loss')
```

```


ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.07039678009253549
Test accuracy: 0.9829



In [25]:

```

w_after = model_relu.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```





3.2)MLP with 2 hidden layers + Relu + AdamOptimizer(WithoutBatchNormalisation,Without dropout)

In [21]:

```
model_relu = Sequential()
#Let's take dimension of first, second,third,fourth and fifth hidden layer 650,550,400,200 and 100
respectively
# If we sample weights from a normal distribution  $N(0,\sigma)$  we satisfy this condition with
 $\sigma = \sqrt{2/(n_i)}$ .
# h1 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.055 \Rightarrow N(0,\sigma) = N(0,0.055)$ 
# h2 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.060 \Rightarrow N(0,\sigma) = N(0,0.060)$ 
# h3 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.070 \Rightarrow N(0,\sigma) = N(0,0.070)$ 
# h4 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.1 \Rightarrow N(0,\sigma) = N(0,0.1)$ 
# h5 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.141 \Rightarrow N(0,\sigma) = N(0,0.141)$ 

model_relu.add(Dense(650, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.055, seed=None)))

model_relu.add(Dense(550, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.060, seed=None)))

model_relu.add(Dense(400, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.070, seed=None)))

model_relu.add(Dense(200, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.1, seed=None)))

model_relu.add(Dense(100, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.141, seed=None)))

model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()
```

Layer (type)	Output Shape	Param #
dense_8 (Dense)	(None, 650)	510250
dense_9 (Dense)	(None, 550)	358050
batch_normalization_4 (Batch Normalization)	(None, 550)	2200
dense_10 (Dense)	(None, 400)	220400
dense_11 (Dense)	(None, 200)	80200
batch_normalization_5 (Batch Normalization)	(None, 200)	800

dense_12 (Dense) (None, 100) 20100

dense_13 (Dense) (None, 10) 1010

Total params: 1,193,010

Trainable params: 1,191,510

Non-trainable params: 1,500

In [22]:

```
#Applying adam optimizer
```

```
model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 21s 344us/step - loss: 0.2085 - acc: 0.9375 - val_loss: 0.1278 - val_acc: 0.9608

Epoch 2/20

60000/60000 [=====] - 20s 330us/step - loss: 0.0881 - acc: 0.9732 - val_loss: 0.0956 - val_acc: 0.9712

Epoch 3/20

60000/60000 [=====] - 20s 330us/step - loss: 0.0612 - acc: 0.9807 - val_loss: 0.0983 - val_acc: 0.9704

Epoch 4/20

60000/60000 [=====] - 20s 333us/step - loss: 0.0505 - acc: 0.9839 - val_loss: 0.1061 - val_acc: 0.9701

Epoch 5/20

60000/60000 [=====] - 23s 386us/step - loss: 0.0462 - acc: 0.9847 - val_loss: 0.1037 - val_acc: 0.9705

Epoch 6/20

60000/60000 [=====] - 27s 446us/step - loss: 0.0383 - acc: 0.9876 - val_loss: 0.0902 - val_acc: 0.9744

Epoch 7/20

60000/60000 [=====] - 28s 460us/step - loss: 0.0342 - acc: 0.9886 - val_loss: 0.0921 - val_acc: 0.9758

Epoch 8/20

60000/60000 [=====] - 28s 463us/step - loss: 0.0309 - acc: 0.9896 - val_loss: 0.0986 - val_acc: 0.9735

Epoch 9/20

60000/60000 [=====] - 33s 546us/step - loss: 0.0261 - acc: 0.9914 - val_loss: 0.1098 - val_acc: 0.9700

Epoch 10/20

60000/60000 [=====] - 34s 567us/step - loss: 0.0225 - acc: 0.9926 - val_loss: 0.0871 - val_acc: 0.9771

Epoch 11/20

60000/60000 [=====] - 30s 506us/step - loss: 0.0206 - acc: 0.9931 - val_loss: 0.0758 - val_acc: 0.9795

Epoch 12/20

60000/60000 [=====] - 31s 512us/step - loss: 0.0235 - acc: 0.9924 - val_loss: 0.1162 - val_acc: 0.9711

Epoch 13/20

60000/60000 [=====] - 28s 464us/step - loss: 0.0230 - acc: 0.9925 - val_loss: 0.0812 - val_acc: 0.9799

Epoch 14/20

60000/60000 [=====] - 28s 472us/step - loss: 0.0164 - acc: 0.9947 - val_loss: 0.1055 - val_acc: 0.9749

Epoch 15/20

60000/60000 [=====] - 28s 460us/step - loss: 0.0195 - acc: 0.9933 - val_loss: 0.1036 - val_acc: 0.9780

Epoch 16/20

60000/60000 [=====] - 28s 467us/step - loss: 0.0179 - acc: 0.9942 - val_loss: 0.0735 - val_acc: 0.9821

Epoch 17/20

60000/60000 [=====] - 27s 452us/step - loss: 0.0147 - acc: 0.9952 - val_loss: 0.0994 - val_acc: 0.9771

Epoch 18/20

60000/60000 [=====] - 29s 481us/step - loss: 0.0122 - acc: 0.9959 - val_loss: 0.0755 - val_acc: 0.9822

Epoch 19/20

60000/60000 [=====] - 28s 473us/step - loss: 0.0128 - acc: 0.9958 - val_loss: 0.0757 - val_acc: 0.9820

Epoch 20/20


```
Epoch 20/20  
60000/60000 [=====] - 28s 470us/step - loss: 0.0143 - acc: 0.9954 - val_loss: 0.0966 - val_acc: 0.9769 - loss: 0.0133 - acc: -
```

In [23]:

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)  
print('Test score:', score[0])  
print('Test accuracy:', score[1])  
  
fig, ax = plt.subplots(1, 1)  
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')  
  
# list of epoch numbers  
x = list(range(1, nb_epoch+1))  
  
vy = history.history['val_loss']  
ty = history.history['loss']  
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.09663725952169334


Test accuracy: 0.9769



In [24]:

```
w_after = model_relu.get_weights()  
  
h1_w = w_after[0].flatten().reshape(-1,1)  
h2_w = w_after[2].flatten().reshape(-1,1)  
out_w = w_after[4].flatten().reshape(-1,1)  
  
fig = plt.figure()  
plt.title("Weight matrices after model trained")  
plt.subplot(1, 3, 1)  
plt.title("Trained model Weights")  
ax = sns.violinplot(y=h1_w, color='b')  
plt.xlabel('Hidden Layer 1')  
  
plt.subplot(1, 3, 2)  
plt.title("Trained model Weights")  
ax = sns.violinplot(y=h2_w, color='r')  
plt.xlabel('Hidden Layer 2')  
  
plt.subplot(1, 3, 3)
```

```
plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



In [32]:

```
x = PrettyTable()
x.field_names = ["No of layer", "Dropout&normalizationused", "Accuracy"]
x.add_row(["2", "Yes", "0.9839"])
x.add_row(["2", "No", "0.9835"])
x.add_row(["2", "Yes", "0.9809"])
x.add_row(["3", "No", "0.9786"])
x.add_row(["3", "Yes", "0.9829"])
x.add_row(["5", "NO", "0.9769"])

print(x)
```

No of layer	Dropout&normalizationused	Accuracy
2	Yes	0.9839
2	No	0.9835
2	Yes	0.9809
3	No	0.9786
3	Yes	0.9829
5	NO	0.9769

Conclusion

By the above table we can infer that adding dropouts and by batchnormalizing, performance of the model increases